

# A Gentle Introduction to Python

**Phuong Vu**

Department of Biostatistics  
UNIVERSITY OF WASHINGTON

December 6, 2017

# Acknowledgment



# What is Python?



- A programming language
- Quite similar to R
- Take spacing **very seriously**

# Why Use Python?

According to Noah Simon:

- Hacking around with a robot
- Building a website
- Good machine learning libraries
- Better/cleaner for building big systems
- Inference is pretty meh...

My reason for learning Python?...

# Why Use Python?

Sometimes Python is just faster than R, especially with loops...

In R

```
ptm1 <- proc.time()
for(i in 1:100000) {
  print("Hi")
}
print("Finished!")
ptm2 <- proc.time()

print(ptm2 - ptm1, digits = 4)
```

Result:

```
user  system elapsed
2.67   0.08    2.63
```

In Python

```
import time

timer_start = time.time()
for i in xrange(1, 100000):
  print "Hi!"
print "Finished!"
timer_end = time.time()

print timer_end - timer_start
```

Result:

```
0.946000099182
```

# How to Run Python?

```
bayes.biostat.washington.edu - phuongvu@bayes0:~ - VT
File Edit Setup Control Window Help
Last login: Mon Oct 30 15:19:58 2017 from d-173-250-178-11.dhcp4.washington.edu
Welcome to bayes.biostat.washington.edu.
NEWS
----
09/06/17 -- R 3.4.1 installed
09/22/17 -- R 3.4.1 is now the default version
            (replacing R 3.3.2)
            To use a previous version of R, launch it by
            its version-specific name, e.g. 'R-3.3.2'

[phuongvu@bayes0 ~]$ python2.7
Python 2.7.11 (default, Feb  4 2016, 12:37:41)
[GCC 4.4.6 20120305 (Red Hat 4.4.6-4)] on linux2
Type "help", "copyright", "credits" or "license()" for more information.
>>>
```



# How to Run Python?

The screenshot displays the Spyder Python IDE interface. The main window is titled "Spyder (Python 2.7)". The menu bar includes File, Edit, Search, Source, Run, Debug, Consoles, Projects, Tools, View, and Help. The toolbar contains icons for file operations, running, and debugging. The editor pane shows a file named "GenerateSurface.py" with the following code:

```
1 #!usr/local/bin/python2.7
2
3 #####
4 ## FILE:   Functions_SpatialCovMat
5 ## CREATED: 20170405
6 ##
7 ## AUTHOR:  Phuong T. Vu
8 ##
9 ## PURPOSE: function to generate spatial covariance matrix with exponential
10 ##
11 ## YYYYYYDD  INIT  COMMENTS
12 ## -----  ---  -----
13 ## 20170405  PTV   Create file (separated from previous file)
14 #####
15
16 ## FUNCTION: GenerateCovMat - generate the V matrix for Latent variable
17 ## ARGS:    s - the coordinates of the locations (should be a n-by-2 matrix)
18 ##          tau2 - nugget
19 ##          sigma2 - partial sill
20 ##          phi - 1/range
21 ## RETURNS: V matrix
22 def GenerateCovMat(s, sigma2, tau2, phi):
23     import numpy as np
24     from scipy.spatial.distance import pdist, squareform
25     n = s.shape[0] ## get the number of locations
26     dist = squareform(pdist(s)) ## compute pairwise distance
27     return sigma2 * np.exp(-dist/phi) + np.diag(np.repeat(tau2, n))
```

The Variable explorer on the right shows a table with columns Name, Type, Size, and Value. Below it, the File explorer is visible. The IPython console at the bottom right shows the following text:

```
Python 2.7.12 [Anaconda custom (64-bit)] (default, Jun 29 2016, 11:07:13) [MSC v.1500 64 bit (AMD64)]
Type "copyright", "credits" or "license" for more information.

IPython 5.1.0 -- An enhanced Interactive Python.
?                -> Introduction and overview of IPython's features.
?quickref        -> Quick reference.
help             -> Python's own help system.
object?         -> Details about 'object', use 'object??' for extra details.

In [1]:
```

The status bar at the bottom indicates: IPython console | History log | Permissions: RW | End-of-lines: CRLF | Encoding: ASCII | Line: 5 | Column: 22 | Memory: 63 %

## Python 2.x or Python 3.x?

- Official on Python Wiki: *Python 2.x is legacy, Python 3.x is the present and future of the language*
- Final version of 2.x is 2.7, and for 3.x is 3.6 released in 2016
- Pros and cons:
  - 3.x is on the cutting edge, i.e. moving forward new features will be implemented on 3.x but not added to 2.x
  - 2.x has much better library support and documentation
  - 2.x has much more extensive and specific third-party packages or utility that may not have been released on 3.x yet
  - Some Linux distributions and Macs still use 2.x as default
- Me? Python 2.7    ͡\_(`͡)\_/͡- it's just a tool...




# Python modules

- To do any serious (statistical) work in Python, you need to extend it with modules, which are similar to R libraries
- Frequently used modules often come with the Anaconda distribution or are already installed on the department clusters: `numpy`, `scipy`, `pandas`, `scikit-learn`
- Other stuff you may like: `CvxPy` (convex optimization), `TensorFlow` by Google (deep learning)

# Python modules

- To do any serious (statistical) work in Python, you need to extend it with modules, which are similar to R libraries
- Frequently used modules often come with the Anaconda distribution or are already installed on the department clusters: `numpy`, `scipy`, `pandas`, `scikit-learn`
- Other stuff you may like: `CvxPy` (convex optimization), TensorFlow by Google (deep learning)
- How much of these cool modules will we cover today? Not much!



## How Similar is Python to R?

- Python is a lot like R - there are lists, strings, arrays, and functions which should all seem similar, although some syntax is different
- Indentation/spacing are taken **very seriously!**
- Indexing starts at **0, not 1!**
- Help files are very well-documented:
  - For python2.x: <https://docs.python.org/2/tutorial/>
  - For python3.x: <https://docs.python.org/3/tutorial/>

# Methods and Objects

- Everything in Python is an object

- Objects have classes:

```
In [1]: x = "statistics"
```

```
In [2]: x.__class__
```

```
Out[2]: str
```

- Each class has methods and attributes associated with objects of that type. Access these with the "." operator. Find out what methods and object has with dir.

```
In [3]: dir(x)
```

```
Out[3]: <omit for space limit>
```

```
In [4]: x.capitalize()
```

```
Out[4]: 'Statistics'
```

```
In [5]: print x.__doc__
```

```
Out[5]: str(object='') -> string
```

Return a nice string representation of the object.

If the argument is a string, the return value is the same object.

- You can get a long way just using the built-in methods

## Example: Loops in Python

Silly example:

```
import numpy as np

x = np.zeros(shape = (11, 2)) ## create a 11x2 array of zeros
x[:, 1] = np.linspace(start = 0, stop = 10, num = 11) ## replace
the 2nd column with some non-zeros

## this is just to show how a for loop in python looks like...
for i in range(11):
    x[i, 0] = x[i, 1]/3 + 1

## you should probably do this instead...
x[:, 0] = x[:, 1]/3 + 1
```

## Example: Matrix Algebra in Python

```
import numpy as np

A = np.array([[1, 2, 3], [4, 5, 6]]) ## create a 2x3 array
A.T ## transpose of A

## this gives you error
b1 = np.array([0.5, 1])
A.dot(b1) ## mismatched shapes

## this runs, but final answer is a (2, ) array
b2 = np.array([0.5, 1, 2])
A.dot(b2) ## shapes are not aligned

## this runs, and final answer is a (2, 1) array (column vector)
A.dot(np.reshape(b2, (3, 1)))
```

- Arrays and matrices are two different objects in python
- You should almost always stick with arrays
- Make sure you always know/check the dimensions

## Example: Data manipulation in Python

```
import numpy as np
import pandas as pd

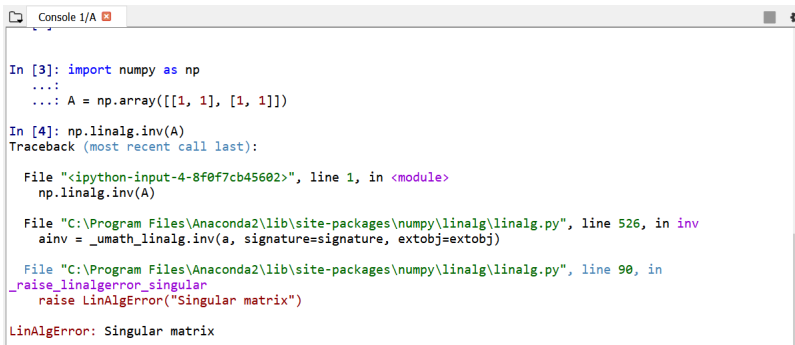
## read a csv file
data = pd.read_csv('exposure_surface.csv')
data['covariate.1'] ## see the variable named "covariate.1"

## convert an array into a pandas dataframe and print to a .csv file
n = 100
p = 5
fakeX = np.reshape(np.random.normal(0, 1, n*p), (n, p))
fakeY = np.random.normal(0, 1, n)

fakedata = np.column_stack( (fakeX, fakeY) )
cols = ['var.' + str(i+1) for i in range(p)] + ['outcome']
fakedata_pd = pd.DataFrame(fakedata, columns = cols)
fakedata_pd.to_csv("fakedata.csv", index = False)
```

# Excellent Error Traceback

Python tells you exactly what and where (i.e. on which line of which function/module) the error is!



```
Console 1/A

In [3]: import numpy as np
...:
...: A = np.array([[1, 1], [1, 1]])

In [4]: np.linalg.inv(A)
Traceback (most recent call last):

  File "<ipython-input-4-8f0f7cb45602>", line 1, in <module>
    np.linalg.inv(A)

  File "C:\Program Files\Anaconda2\lib\site-packages\numpy\linalg\linalg.py", line 526, in inv
    ainv = _umath_linalg.inv(a, signature=signature, extobj=extobj)

  File "C:\Program Files\Anaconda2\lib\site-packages\numpy\linalg\linalg.py", line 90, in
_raise_linalgerror_singular
    raise LinAlgError("Singular matrix")

LinAlgError: Singular matrix
```



## Alias - Friend or Foe?

- When two identifiers refer to the same variable (and therefore value), this is known as an **alias**
- In R, `X <- f(X)` generally makes a copy of X, which sometimes mean slower performance due to making too many copies of data... In Python, `f(X)` needs not...

## Alias - Friend or Foe?

Sometimes aliasing is bad... For example:

- In R, you sometimes do this to generate placeholders:

```
>>> output1 <- output2 <- rep(0, nsim)
```

and changing output2 does not change output1

- In Python, you will end up with 2 output arrays that look exactly the same...

```
In [1]: import numpy as np
```

```
In [2]: output1 = output2 = np.zeros((5))
```

```
In [3]: print id(output1), id(output2)
```

```
Out[3]: 204788400 204788400
```

```
In [4]: output1[1] = 0.5 ## change output1
```

```
In [5]: output2 ## check output2...
```

```
Out[5]: array([ 0. ,  0.5,  0. ,  0. ,  0. ])
```

## Run as Script?

- Similar to R, you will need this line for the .py file to be executed on Cox or Bayes/Gosset:

```
#!/usr/local/bin/python2.7  
<your python code>
```

- Submit on Cox:

```
[phuongvu@cox ~]$ /usr/local/bin/python2.7 Sim.py &
```

## Run as Script?

- Submit on Bayes/Gosset:

```
[phuongvu@bayes0 Dissertation]$ ./submit_Sim.sh
```

```
[phuongvu@bayes0 Dissertation]$ cat submit_Sim.sh
```

```
#!/bin/sh
```

```
qsub -t 1-22 -cwd -e Trashfiles/ -o Trashfiles/ -q s-normal.q  
-M phuongvu@uw.edu -m e call_Sim.sh
```

```
[phuongvu@bayes0 Dissertation]$ cat call_Sim.sh
```

```
#!/bin/sh
```

```
<extra stuff omitted>
```

```
/usr/local/bin/python2.7 Sim.py
```

- The current underlying BLAS libraries for python2.7 on Bayes/Gosset are NOT optimized. If you seriously consider python for your research, you should consider link your modules with a more optimized BLAS library (ask BITE!)

## Final Remarks

- The best way to learn Python is diving into some coding projects and using google/stackoverflow/etc.
- Python2.x or 3.x? Just pick one!
- OOP, by definition, is designed in a manner to minimize the amount of new code that needs to be written. There maybe a well-written and well-tested module somewhere that can do the job you want.
- For those who are interested in spatial statistics or GIS stuff: Python is the primary scripting language for ArcGIS
- Don't hesitate to ask your fellow students for tips and tricks!

## Homework 7\\_(ツ)\_/

Write a script in Python to execute the following:

1. Choose a vector  $\beta \in \mathbb{R}^p$  for  $p = 3$ , and generate data:

$$\mathbf{y} = \mathbf{X}\beta + \epsilon$$

$$\mathbf{X} \in \mathbb{R}^{n \times p}, \mathbf{y} \in \mathbb{R}^n, \epsilon \in \mathbb{R}^n, n = 300$$

$$\mathbf{X}_{i1} = 1 \text{ for } i = 1, \dots, n$$

$$\mathbf{X}_{i2} \sim \text{Bernoulli}(0.7) \text{ for } i = 1, \dots, n$$

$$\mathbf{X}_{i3} \sim \text{Uniform}(-1, 1) \text{ for } i = 1, \dots, n$$

$$\epsilon_i \sim \mathcal{N}(0, 1) \text{ for } i = 1, \dots, n$$

2. Find  $\hat{\beta} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$
3. Save your generated  $\mathbf{X}$  and  $\mathbf{y}$  to a .csv file

In R, read in this .csv file, and confirm that your result, either by doing the algebra or using `lm()`

You will only need to google functions from `numpy` and `pandas`.